

# Self-Driving Car Engineer Nanodegree

## Deep Learning

### Project: Build a Traffic Sign Recognition Classifier

In this notebook, a template is provided for you to implement your functionality in stages which is required to successfully complete this project. If additional code is required that cannot be included in the notebook, be sure that the Python code is successfully imported and included in your submission, if necessary. Sections that begin with '**Implementation**' in the header indicate where you should begin your implementation for your project. Note that some sections of implementation are optional, and will be marked with '**Optional**' in the header.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

---

## Step 1: Dataset Exploration

Visualize the German Traffic Signs Dataset. This is open ended, some suggestions include: plotting traffic signs images, plotting the count of each sign, etc. Be creative!

The pickled data is a dictionary with 4 key/value pairs:

- features -> the images pixel values, (width, height, channels)
- labels -> the label of the traffic sign
- sizes -> the original width and height of the image, (width, height)
- coords -> coordinates of a bounding box around the sign in the image, (x1, y1, x2, y2). Based the original image (not the resized version).

In [1]:

```
# Load pickled data
import pickle

# TODO: fill this in based on where you saved the training and testing data
training_file = 'traffic-signs-data/train.p'
testing_file = 'traffic-signs-data/test.p'

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_test, y_test = test['features'], test['labels']
```

In [2]:

```
### To start off let's do a basic data summary.

# TODO: number of training examples
n_train = len(X_train)

# TODO: number of testing examples
n_test = len(X_test)

# TODO: what's the shape of an image?
image_shape = X_train[0].shape

# TODO: how many classes are in the dataset
n_classes = len(set(y_train))

print("Number of training examples =", n_train)
print("Number of testing examples =", n_test)
print("Image data shape =", image_shape)
print("Number of classes =", n_classes)
```

```
Number of training examples = 39209
Number of testing examples = 12630
Image data shape = (32, 32, 3)
Number of classes = 43
```

In [3]:

```
### Data exploration visualization goes here.
### Feel free to use as many code cells as needed.

import matplotlib.pyplot as plt
%matplotlib inline

print("Original image shape =", train["sizes"].shape)
print("Image coordinate shape:", train["coords"].shape)

#plot image helper function
def plot_images(images, labels, cls_pred=None):
    # Plot image.
    plt.imshow(images, cmap='binary')


    # Show true and predicted classes.
    if cls_pred is None:
        xlabel = "Label: {0}".format(labels)
    else:
        xlabel = "Label: {0}, Pred: {1}".format(labels, cls_pred)

    plt.xlabel(str(xlabel))

    # Remove ticks from the plot.
    #     plt.xticks([])
    #     plt.yticks([])

plot_images(X_train[0], y_train[0])
```

Original image shape = (39209, 2)  
Image coordinate shape: (39209, 4)



---

## Step 2: Design and Test a Model Architecture

Design and implement a deep learning model that learns to recognize traffic signs. Train and test your model on the [German Traffic Sign Dataset](http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset) (<http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset>).

There are various aspects to consider when thinking about this problem:

- Your model can be derived from a deep feedforward net or a deep convolutional network.
- Play around preprocessing techniques (normalization, rgb to grayscale, etc)
- Number of examples per label (some have more than others).
- Generate fake data.

Here is an example of a [published baseline model on this problem](http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf) (<http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf>). It's not required to be familiar with the approach used in the paper but, it's good practice to try to read papers like these.

# Implementation

Use the code cell (or multiple code cells, if necessary) to implement the first step of your project. Once you have completed your implementation and are satisfied with the results, be sure to thoroughly answer the questions that follow.

In [4]:

```
### Preprocess the data here.
### Feel free to use as many code cells as needed.

import numpy as np
import pandas as pd
import cv2

#convert RGB to gray
features_train = []
features_test = []

for i, each in enumerate(X_train):
    each = cv2.cvtColor(each, cv2.COLOR_BGR2GRAY)
    each = each.reshape([32, 32, 1])
    features_train.append(each)
features_train = np.asarray(features_train)/255

for i, each in enumerate(X_test):
    each = cv2.cvtColor(each, cv2.COLOR_BGR2GRAY)
    each = each.reshape([32, 32, 1])
    features_test.append(each)
features_test = np.asarray(features_test)/255

# read and count number of sign names for each class
label_names = pd.read_csv("signnames.csv")
label_names.head(5)
```

Out[4]:

	ClassId	SignName
0	0	Speed limit (20km/h)
1	1	Speed limit (30km/h)
2	2	Speed limit (50km/h)
3	3	Speed limit (60km/h)
4	4	Speed limit (70km/h)

In [5]:

```
# count number of signs for each sign name
```

## Question 1

*Describe the techniques used to preprocess the data.*

**Answer:** Images are first converted from RGB 3 channels to gray scale one channel, then images are divided by 255 to normalize between 0 and 1. The size of images are also reshaped to (32, 32, 1) for easy processing with tensorflow. Data set features were split into training and validation set with validation set size at least 3000 to ensure a 1% true accuracy. Data set labels were one-hot encoded using sklearn LabelBinarizer method.

In [6]:

```
### Generate data additional (if you want to!)
### and split the data into training/validation/testing sets here.
### Feel free to use as many code cells as needed.

# Turn labels into numbers and apply One-Hot Encoding
from sklearn.preprocessing import LabelBinarizer

encoder = LabelBinarizer()
encoder.fit(y_train)
labels_train = encoder.transform(y_train)
encoder = LabelBinarizer()
encoder.fit(y_test)
labels_test = encoder.transform(y_test)

# Change to float32, so it can be multiplied against the features in TensorFlow,
which are float32
labels_train = labels_train.astype(np.float32)
labels_test = labels_test.astype(np.float32)
is_labels_encod = True

print('Labels One-Hot Encoded')
print(labels_train.shape)
print(labels_test.shape)
```

```
Labels One-Hot Encoded
(39209, 43)
(12630, 43)
```

## Question 2

*Describe how you set up the training, validation and testing data for your model. If you generated additional data, why?*

**Answer:** Validation data is extraced from training data using sklearn train\_test\_split method with test\_size 20 percent of original training data to ensure validation data size is at least 3000. No additional data is generated below with this commit, maybe rotating, cropping the image will provide more data to gain better accuracy

In [7]:

```
### Define your architecture here.  
### Feel free to use as many code cells as needed.
```

In [8]:

```
# split training and validation data  
# Get randomized datasets for training and validation  
# Validation data set at least 3000 to guarantee a 1% true accuracy  
from sklearn.model_selection import train_test_split  
  
features_train, features_valid, labels_train, labels_valid = train_test_split(  
    features_train,  
    labels_train,  
    test_size=0.2)  
  
print('Training features and labels randomized and split.')
```

```
print("Training set features and labels size:", features_train.shape, labels_train.shape)  
print("Validation set features and labels size:", features_valid.shape, labels_valid.shape)
```

```
Training features and labels randomized and split.  
Training set features and labels size: (31367, 32, 32, 1) (31367, 43)  
)  
Validation set features and labels size: (7842, 32, 32, 1) (7842, 43)  
)
```

In [9]:

```
# 2-layer CNN followed by 2 fully-connected layer  
  
# import libraries  
import time  
from datetime import timedelta  
import math  
import tensorflow as tf  
  
# configurations of CNN  
# Convolutional Layer 1.  
filter_size1 = 5          # Convolution filters are 5 x 5 pixels.  
num_filters1 = 16         # There are 16 of these filters.  
  
# Convolutional Layer 2.  
filter_size2 = 5          # Convolution filters are 5 x 5 pixels.  
num_filters2 = 36         # There are 36 of these filters.
```

```

# Fully-connected layer.
fc_size = 128          # Number of neurons in fully-connected layer.

# data dimensions
img_size = 32
img_size_flat = img_size * img_size
img_shape = (img_size, img_size)
num_channels = 1
num_classes = 43

# helper-functions
def new_weights(shape):
    return tf.Variable(tf.truncated_normal(shape, stddev=0.05))

def new_biases(length):
    return tf.Variable(tf.constant(0.05, shape=[length]))

def new_conv_layer(input, num_input_channels, filter_size, num_filters, keep_prob, use_pooling=True):
    shape = [filter_size, filter_size, num_input_channels, num_filters]
    weights = new_weights(shape)
    biases = new_biases(length=num_filters)
    layer = tf.nn.conv2d(input=input, filter=weights, strides=[1,1,1,1], padding='SAME')
    layer += biases

    if use_pooling:
        layer = tf.nn.max_pool(value=layer, ksize=[1,2,2,1], strides=[1,2,2,1], padding='SAME')

    layer = tf.nn.relu(layer)
    layer = tf.nn.dropout(layer, keep_prob)
    return layer, weights

def flatten_layer(layer):
    layer_shape = layer.get_shape()
    num_features = layer_shape[1:4].num_elements()
    layer_flat = tf.reshape(layer, [-1, num_features])
    return layer_flat, num_features

def new_fc_layer(input, num_inputs, num_outputs, use_relu=True):
    weights = new_weights(shape=[num_inputs, num_outputs])
    biases = new_biases(length=num_outputs)
    layer = tf.matmul(input, weights) + biases

    if use_relu:
        layer = tf.nn.relu(layer)
    return layer

```



### Question 3

*What does your final architecture look like? (Type of model, layers, sizes, connectivity, etc.) For reference on how to build a deep neural network using TensorFlow, see [Deep Neural Network in TensorFlow](https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/b516a270-8600-4f93-a0a3-20dfeabe5da6/concepts/83a3a2a2-a9bd-4b7b-95b0-eb924ab14432) (<https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/b516a270-8600-4f93-a0a3-20dfeabe5da6/concepts/83a3a2a2-a9bd-4b7b-95b0-eb924ab14432>) from the classroom.*

**Answer:** Flow graph is defined as below. It's a basic 2-layer CNN with keep\_prob rate equals to 0.7 followed by 2 fully connected layers. The filter size for both CNN is 5x5. The first convnet channel is 16, second convnet channel is 36, first fully connected layer has 128 neurons. Both convnets use same padding, have strides equal to (1,1,1,1) and use max pooling with strides equal to (1,2,2,1).

In [10]:

```
### Train your model here.
### Feel free to use as many code cells as needed.

## Tensor flow graph

# placeholder variables

x = tf.placeholder(tf.float32, shape=[None, img_size, img_size, num_channels])
y_true = tf.placeholder(tf.float32, shape=[None, num_classes], name='y_true')
y_true_cls = tf.argmax(y_true, dimension=1)

# Convnet 1
convnet1, weights1 = new_conv_layer(input=x, num_input_channels=num_channels, filter_size=filter_size1, num_filters=num_filters1, keep_prob=0.7, use_pooling=True)

# Convnet 2
convnet2, weights2 = new_conv_layer(input=convnet1, num_input_channels=num_filters1, filter_size=filter_size2, num_filters=num_filters2, keep_prob=0.7, use_pooling=True)

# Flatten convnet 2
layer_flat, num_features = flatten_layer(convnet2)

# FC1
fc1 = new_fc_layer(input=layer_flat, num_inputs=num_features, num_outputs=fc_size, use_relu=False)

# FC2
fc2 = new_fc_layer(input=fc1, num_inputs=fc_size, num_outputs=num_classes, use_relu=False)

# predicted label
y_pred = tf.nn.softmax(fc2)
y_pred_cls = tf.argmax(y_pred, 1)

print("Convnet 1 shape:", convnet1)
print("Convnet 2 shape:", convnet2)
print("Flattened Convnet2 shape:", layer_flat)
print("FC1 shape:", fc1)
print("FC2 shape:", fc2)
```

```
Convnet 1 shape: Tensor("dropout/mul:0", shape=(?, 16, 16, 16), dtype=float32)
Convnet 2 shape: Tensor("dropout_1/mul:0", shape=(?, 8, 8, 36), dtype=float32)
Flattened Convnet2 shape: Tensor("Reshape:0", shape=(?, 2304), dtype=float32)
FC1 shape: Tensor("add_2:0", shape=(?, 128), dtype=float32)
FC2 shape: Tensor("add_3:0", shape=(?, 43), dtype=float32)
```

## Question 4

*How did you train your model? (Type of optimizer, batch size, epochs, hyperparameters, etc.)*

**Answer:** AdamOptimizer, an advanced gradient descent optimizer is used with learning rate equal to 0.0001. 100 iterations are made with batch size equal to 128 in each iteration.

In [11]:

```
# optimization
learning_rate = 0.0001
epoches = 100
batch_size = 128

cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits=fc2, labels=y_true)
cost = tf.reduce_mean(cross_entropy)
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost)
correct_prediction = tf.equal(y_pred_cls, y_true_cls)
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

In [12]:

```
# tensorflow run
# open and init session
session = tf.Session()
session.run(tf.initialize_all_variables())

# randomly select batches from train data
def random_batch(features, labels):
    # Number of images in the training-set.
    num_images = len(features)

    # Create a random index.
    idx = np.random.choice(num_images,
                           size=batch_size,
                           replace=False)

    # Use the random index to select random images and labels.
    features_batch = features[idx, ]
    labels_batch = labels[idx, ]
    return features_batch, labels_batch

feed_dict_validation = {x: features_valid, y_true: labels_valid}
feed_dict_test = {x: features_test, y_true: labels_test}

for i in range(epochs):
    start_time = time.time()
    total_batch = int(features_train.shape[0]/batch_size)
    for j in range(total_batch):
        batch_x, batch_y = random_batch(features_train, labels_train)
        feed_dict_train = {x: batch_x, y_true: batch_y}
        session.run(optimizer, feed_dict=feed_dict_train)
        acc_train = session.run(accuracy, feed_dict=feed_dict_train)

    acc_validation = session.run(accuracy, feed_dict=feed_dict_validation)
    acc_test = session.run(accuracy, feed_dict =feed_dict_test)
    end_time = time.time()
    time_diff = end_time - start_time
    print("Epoch:", i)
    print("Time usage: " + str(timedelta(seconds=int(round(time_diff)))))
    print("training accuracy is:", acc_train)
    print("validation set accuracy is:", acc_validation)
    print("test set accuracy is:", acc_test)
```

```
Epoch: 0
Time usage: 0:02:02
training accuracy is: 0.0859375
validation set accuracy is: 0.0992094
test set accuracy is: 0.0976247
Epoch: 1
Time usage: 0:01:53
training accuracy is: 0.21875
validation set accuracy is: 0.1696
```

test set accuracy is: 0.14426  
Epoch: 2  
Time usage: 0:01:50  
training accuracy is: 0.265625  
validation set accuracy is: 0.293548  
test set accuracy is: 0.257403  
Epoch: 3  
Time usage: 0:01:49  
training accuracy is: 0.46875  
validation set accuracy is: 0.423106  
test set accuracy is: 0.372763  
Epoch: 4  
Time usage: 0:01:49  
training accuracy is: 0.46875  
validation set accuracy is: 0.514027  
test set accuracy is: 0.464529  
Epoch: 5  
Time usage: 0:01:49  
training accuracy is: 0.632812  
validation set accuracy is: 0.590538  
test set accuracy is: 0.521378  
Epoch: 6  
Time usage: 0:01:50  
training accuracy is: 0.664062  
validation set accuracy is: 0.649197  
test set accuracy is: 0.574426  
Epoch: 7  
Time usage: 0:01:48  
training accuracy is: 0.695312  
validation set accuracy is: 0.683117  
test set accuracy is: 0.605701  
Epoch: 8  
Time usage: 0:01:48  
training accuracy is: 0.765625  
validation set accuracy is: 0.714996  
test set accuracy is: 0.641964  
Epoch: 9  
Time usage: 0:01:48  
training accuracy is: 0.703125  
validation set accuracy is: 0.748789  
test set accuracy is: 0.668171  
Epoch: 10  
Time usage: 0:01:49  
training accuracy is: 0.78125  
validation set accuracy is: 0.768426  
test set accuracy is: 0.682027  
Epoch: 11  
Time usage: 0:01:48  
training accuracy is: 0.820312  
validation set accuracy is: 0.787299  
test set accuracy is: 0.701425  
Epoch: 12  
Time usage: 0:01:48

training accuracy is: 0.710938  
validation set accuracy is: 0.804259  
test set accuracy is: 0.715439  
Epoch: 13  
Time usage: 0:01:48  
training accuracy is: 0.851562  
validation set accuracy is: 0.816501  
test set accuracy is: 0.732621  
Epoch: 14  
Time usage: 0:01:48  
training accuracy is: 0.890625  
validation set accuracy is: 0.827085  
test set accuracy is: 0.74228  
Epoch: 15  
Time usage: 0:01:48  
training accuracy is: 0.867188  
validation set accuracy is: 0.835119  
test set accuracy is: 0.749644  
Epoch: 16  
Time usage: 0:01:51  
training accuracy is: 0.875  
validation set accuracy is: 0.846595  
test set accuracy is: 0.763579  
Epoch: 17  
Time usage: 0:01:49  
training accuracy is: 0.882812  
validation set accuracy is: 0.852461  
test set accuracy is: 0.776564  
Epoch: 18  
Time usage: 0:01:51  
training accuracy is: 0.875  
validation set accuracy is: 0.858327  
test set accuracy is: 0.777514  
Epoch: 19  
Time usage: 0:01:48  
training accuracy is: 0.875  
validation set accuracy is: 0.867508  
test set accuracy is: 0.781473  
Epoch: 20  
Time usage: 0:01:53  
training accuracy is: 0.882812  
validation set accuracy is: 0.8772  
test set accuracy is: 0.799604  
Epoch: 21  
Time usage: 0:01:48  
training accuracy is: 0.914062  
validation set accuracy is: 0.880515  
test set accuracy is: 0.801029  
Epoch: 22  
Time usage: 0:01:48  
training accuracy is: 0.921875  
validation set accuracy is: 0.888931  
test set accuracy is: 0.807997

Epoch: 23  
Time usage: 0:01:48  
training accuracy is: 0.929688  
validation set accuracy is: 0.888931  
test set accuracy is: 0.802613  
Epoch: 24  
Time usage: 0:01:48  
training accuracy is: 0.875  
validation set accuracy is: 0.89875  
test set accuracy is: 0.813935  
Epoch: 25  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.901811  
test set accuracy is: 0.816785  
Epoch: 26  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.904361  
test set accuracy is: 0.823832  
Epoch: 27  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.904106  
test set accuracy is: 0.825178  
Epoch: 28  
Time usage: 0:01:48  
training accuracy is: 0.890625  
validation set accuracy is: 0.910099  
test set accuracy is: 0.832858  
Epoch: 29  
Time usage: 0:01:49  
training accuracy is: 0.960938  
validation set accuracy is: 0.914945  
test set accuracy is: 0.836421  
Epoch: 30  
Time usage: 0:01:48  
training accuracy is: 0.945312  
validation set accuracy is: 0.920811  
test set accuracy is: 0.834521  
Epoch: 31  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.923234  
test set accuracy is: 0.841409  
Epoch: 32  
Time usage: 0:01:48  
training accuracy is: 0.9375  
validation set accuracy is: 0.923871  
test set accuracy is: 0.842518  
Epoch: 33  
Time usage: 0:01:48  
training accuracy is: 0.96875

validation set accuracy is: 0.927314  
test set accuracy is: 0.846397  
Epoch: 34  
Time usage: 0:01:48  
training accuracy is: 0.960938  
validation set accuracy is: 0.92961  
test set accuracy is: 0.84616  
Epoch: 35  
Time usage: 0:01:48  
training accuracy is: 0.929688  
validation set accuracy is: 0.927825  
test set accuracy is: 0.845764  
Epoch: 36  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.935476  
test set accuracy is: 0.851702  
Epoch: 37  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.934328  
test set accuracy is: 0.856374  
Epoch: 38  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.937516  
test set accuracy is: 0.855344  
Epoch: 39  
Time usage: 0:01:48  
training accuracy is: 0.960938  
validation set accuracy is: 0.932925  
test set accuracy is: 0.856374  
Epoch: 40  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.942234  
test set accuracy is: 0.85772  
Epoch: 41  
Time usage: 0:01:47  
training accuracy is: 0.945312  
validation set accuracy is: 0.937388  
test set accuracy is: 0.859937  
Epoch: 42  
Time usage: 0:01:48  
training accuracy is: 0.960938  
validation set accuracy is: 0.943382  
test set accuracy is: 0.861045  
Epoch: 43  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.94657  
test set accuracy is: 0.864608  
Epoch: 44



Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.941724  
test set accuracy is: 0.862391  
Epoch: 45  
Time usage: 0:01:49  
training accuracy is: 0.945312  
validation set accuracy is: 0.942617  
test set accuracy is: 0.865479  
Epoch: 46  
Time usage: 0:01:47  
training accuracy is: 0.9375  
validation set accuracy is: 0.949503  
test set accuracy is: 0.861362  
Epoch: 47  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.947845  
test set accuracy is: 0.865717  
Epoch: 48  
Time usage: 0:01:48  
training accuracy is: 0.960938  
validation set accuracy is: 0.948483  
test set accuracy is: 0.872051  
Epoch: 49  
Time usage: 0:01:49  
training accuracy is: 0.984375  
validation set accuracy is: 0.950778  
test set accuracy is: 0.868804  
Epoch: 50  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.951543  
test set accuracy is: 0.870467  
Epoch: 51  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.954221  
test set accuracy is: 0.868884  
Epoch: 52  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.954093  
test set accuracy is: 0.87403  
Epoch: 53  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.954603  
test set accuracy is: 0.874426  
Epoch: 54  
Time usage: 0:01:48  
training accuracy is: 0.960938  
validation set accuracy is: 0.958046

test set accuracy is: 0.874426  
Epoch: 55  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.955496  
test set accuracy is: 0.875772  
Epoch: 56  
Time usage: 0:01:47  
training accuracy is: 0.960938  
validation set accuracy is: 0.957026  
test set accuracy is: 0.876722  
Epoch: 57  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.962127  
test set accuracy is: 0.879335  
Epoch: 58  
Time usage: 0:01:49  
training accuracy is: 0.984375  
validation set accuracy is: 0.959322  
test set accuracy is: 0.87601  
Epoch: 59  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.959194  
test set accuracy is: 0.881314  
Epoch: 60  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.957664  
test set accuracy is: 0.878781  
Epoch: 61  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.962127  
test set accuracy is: 0.882027  
Epoch: 62  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.961872  
test set accuracy is: 0.880998  
Epoch: 63  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.961489  
test set accuracy is: 0.883848  
Epoch: 64  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.961234  
test set accuracy is: 0.880602  
Epoch: 65  
Time usage: 0:01:48

training accuracy is: 0.992188  
validation set accuracy is: 0.959959  
test set accuracy is: 0.88171  
Epoch: 66  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.96608  
test set accuracy is: 0.885748  
Epoch: 67  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.963785  
test set accuracy is: 0.885194  
Epoch: 68  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.967355  
test set accuracy is: 0.88654  
Epoch: 69  
Time usage: 0:01:48  
training accuracy is: 0.960938  
validation set accuracy is: 0.964805  
test set accuracy is: 0.883927  
Epoch: 70  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.964677  
test set accuracy is: 0.884086  
Epoch: 71  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.965442  
test set accuracy is: 0.889074  
Epoch: 72  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.962127  
test set accuracy is: 0.888599  
Epoch: 73  
Time usage: 0:01:48  
training accuracy is: 0.953125  
validation set accuracy is: 0.967865  
test set accuracy is: 0.888599  
Epoch: 74  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.968248  
test set accuracy is: 0.889707  
Epoch: 75  
Time usage: 0:01:48  
training accuracy is: 1.0  
validation set accuracy is: 0.967228  
test set accuracy is: 0.885986

Epoch: 76  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.964932  
test set accuracy is: 0.886778  
Epoch: 77  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.969013  
test set accuracy is: 0.890024  
Epoch: 78  
Time usage: 0:01:49  
training accuracy is: 0.984375  
validation set accuracy is: 0.968758  
test set accuracy is: 0.890736  
Epoch: 79  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.968885  
test set accuracy is: 0.888599  
Epoch: 80  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.968885  
test set accuracy is: 0.887569  
Epoch: 81  
Time usage: 0:01:48  
training accuracy is: 0.976562  
validation set accuracy is: 0.969013  
test set accuracy is: 0.894695  
Epoch: 82  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.969268  
test set accuracy is: 0.892082  
Epoch: 83  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.970033  
test set accuracy is: 0.889232  
Epoch: 84  
Time usage: 0:01:48  
training accuracy is: 0.96875  
validation set accuracy is: 0.972584  
test set accuracy is: 0.890499  
Epoch: 85  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.96812  
test set accuracy is: 0.89327  
Epoch: 86  
Time usage: 0:01:48  
training accuracy is: 0.984375

validation set accuracy is: 0.972711  
test set accuracy is: 0.894141  
Epoch: 87  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.972073  
test set accuracy is: 0.892557  
Epoch: 88  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.971563  
test set accuracy is: 0.895012  
Epoch: 89  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.971181  
test set accuracy is: 0.898575  
Epoch: 90  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.971436  
test set accuracy is: 0.893745  
Epoch: 91  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.973604  
test set accuracy is: 0.898575  
Epoch: 92  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.973094  
test set accuracy is: 0.895566  
Epoch: 93  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.975771  
test set accuracy is: 0.896041  
Epoch: 94  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.973221  
test set accuracy is: 0.895804  
Epoch: 95  
Time usage: 0:01:48  
training accuracy is: 0.992188  
validation set accuracy is: 0.972456  
test set accuracy is: 0.896675  
Epoch: 96  
Time usage: 0:01:48  
training accuracy is: 0.984375  
validation set accuracy is: 0.976027  
test set accuracy is: 0.899208  
Epoch: 97

```
Time usage: 0:01:48
training accuracy is: 0.984375
validation set accuracy is: 0.975006
test set accuracy is: 0.898021
Epoch: 98
Time usage: 0:01:48
training accuracy is: 1.0
validation set accuracy is: 0.975899
test set accuracy is: 0.899446
Epoch: 99
Time usage: 0:01:48
training accuracy is: 0.992188
validation set accuracy is: 0.972584
test set accuracy is: 0.896754
```

## Question 5

*What approach did you take in coming up with a solution to this problem?*

**Answer:** I tried a basic linear neural network with a relu first, which yields training and validation accuracy less than 10%. Second try, I added another hidden layer but still accuracy is low less than 50%. Third try, I followed a youtube tensorflow tutorial and came up with this 2 layer CNN, i ran the network overnight and obtained the accuracy above.

---

## Step 3: Test a Model on New Images

Take several pictures of traffic signs that you find on the web or around you (at least five), and run them through your classifier on your computer to produce example results. The classifier might not recognize some local signs but it could prove interesting nonetheless.

You may find `signnames.csv` useful as it contains mappings from the class id (integer) to the actual sign name.

## Implementation


Use the code cell (or multiple code cells, if necessary) to implement the first step of your project. Once you have completed your implementation and are satisfied with the results, be sure to thoroughly answer the questions that follow.



In [61]:

```
### Load the images and plot them here.
### Feel free to use as many code cells as needed.
import matplotlib.image as mpimg
import cv2



img1 = mpimg.imread("traffic-signs-data/60kmh.jpg")
img2 = mpimg.imread("traffic-signs-data/80kmh.jpg")
img3 = mpimg.imread("traffic-signs-data/stop.jpg")
img4 = mpimg.imread("traffic-signs-data/yield.jpg")
img5 = mpimg.imread("traffic-signs-data/no entry.jpg")
img = np.array([img1, img2, img3, img4, img5])
label = ["60kmh", "80kmh", "stop", "yield", "no entry"]

for i, each in enumerate(img):
    each = cv2.cvtColor(each, cv2.COLOR_BGR2GRAY)
    each = cv2.resize(each, (32, 32))
    plt.imshow(each, cmap='gray')
    plt.show()
    img[i] = each.reshape([1, 32, 32, 1])
print(img[0].shape)
```









(1, 32, 32, 1)

### Question 6

Choose five candidate images of traffic signs and provide them in the report. Are there any particular qualities of the image(s) that might make classification difficult? It would be helpful to plot the images in the notebook.

**Answer:** These 5 pictures are downloaded online. A blurrier image should make the classification difficult, also the training images seem all taken on a road with an angle, a normal image facing upfront may be somehow difficult to classify since no such training data provided which also suggests that more data should be generated from image augmentation as suggested in question 2.

In [111]:

```
### Run the predictions here.
### Feel free to use as many code cells as needed.

for i in range(5):
    classification = session.run(y_pred_cls, feed_dict={x: img[i]})
    predicted_label = label_names[label_names["ClassId"]==int(classification)][ "
    SignName"].values
    true_label = label[i]
    print("True label:\t\t", true_label)
    print("Predicted label:\t", predicted_label)
```

```
True label:          60kmh
Predicted label:     ['No passing']
True label:          80kmh
Predicted label:     ['Speed limit (80km/h)']
True label:          stop
Predicted label:     ['Stop']
True label:          yield
Predicted label:     ['Yield']
True label:          no entry
Predicted label:     ['No entry']
```

## Question 7

*Is your model able to perform equally well on captured pictures or a live camera stream when compared to testing on the dataset?*

**Answer:** 4 out of 5 are predicted correctly considering the test accuracy of 0.89 with the testing set, I think this result makes sense. A live camera stream will probably be more accurate.


In [115]:

```
### Visualize the softmax probabilities here.
### Feel free to use as many code cells as needed.


for i in range(5):
    values, indices = session.run(tf.nn.top_k(tf.nn.softmax(y_pred.eval(feed_dict={x: img[i]}), session=session)), k=6))
    print("sign {0} top 6 softmax probabilities are \n {1}\n with indices {2}".format(label[i], values, indices))
    print("\n")
    plt.hist(values)
    plt.show()
```

```
sign 60kmh top 6 softmax probabilities are
[[ 0.06078678  0.02236221  0.02236221  0.02236221  0.02236221  0.02236221]]
with indices [[9 0 1 2 3 4]]
```


```
/Users/yifei/anaconda3/envs/tensorflow/lib/python3.5/site-packages/matplotlib/axes/_axes.py:5882: UserWarning: 2D hist input should be n
samples x nvariables;
this looks transposed (shape is 1 x 6)
'(shape is %d x %d)' % inp.shape[::-1])
```




sign 80kmh top 6 softmax probabilities are  
[[ 0.06078678 0.02236221 0.02236221 0.02236221 0.02236221 0.02236221]]  
with indices [[14 0 1 2 3 4]]



sign stop top 6 softmax probabilities are  
[[ 0.06078678 0.02236221 0.02236221 0.02236221 0.02236221 0.02236221]]  
with indices [[17 0 1 2 3 4]]




sign yield top 6 softmax probabilities are  
[[ 0.06078678 0.02236221 0.02236221 0.02236221 0.02236221 0.02236221]]  
with indices [[13 0 1 2 3 4]]



sign no entry top 6 softmax probabilities are

```
[[ 0.06078678  0.02236221  0.02236221  0.02236221  0.02236221  0.02236221]]
```

with indices `[[17 0 1 2 3 4]]`



## Question 8

Use the model's softmax probabilities to visualize the **certainty** of its predictions, `tf.nn.top_k` ([https://www.tensorflow.org/versions/r0.11/api\\_docs/python/nn.html#top\\_k](https://www.tensorflow.org/versions/r0.11/api_docs/python/nn.html#top_k)) could prove helpful here. Which predictions is the model certain of? Uncertain? If the model was incorrect in its initial prediction, does the correct prediction appear in the top k? (k should be 5 at most)

**Answer:** I printed out top 6 softmax probabilities for each image. First one is predicted wrong but correct prediction does appear at position 5. Other 4 all predicted correctly and the probabilities for correct prediction are all around 6%.

## Question 9

If necessary, provide documentation for how an interface was built for your model to load and classify newly-acquired images.

**Answer:** Images are first read into a numpy array and then converted from RGB to grayscale using cv2 module. Same are then reshaped to size (1, 32, 32, 1) which will be fed to the tensorflow graph. Predictions are made by running the prediction flow.

**Note:** Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

In [ ]: